

Scaling Laws for Decoder-Only Transformers Trained on SVG Code

Aniruth R.
CS-GY 6923 Machine Learning, NYU Tandon
Spring 2026

May 1, 2026

Abstract

We empirically study scaling laws for decoder-only Transformer language models trained on Scalable Vector Graphics (SVG) source code. Across five model sizes from approximately 1M to 88M non-embedding parameters, we fit a Kaplan-style power law $L(N) = a \cdot N^{-\alpha} + c$ to validation loss under standard parameterization (SP), then re-run the sweep under maximal-update parameterization (μ P) and extrapolate the optimal learning rate roughly $10\times$ beyond the largest fitted model. We additionally train the compute-optimal configuration to convergence and report perplexity along with XML well-formedness and render-validity rates on held-out samples.

1 Introduction

Neural scaling laws [?] predict test loss as a smooth power-law function of non-embedding parameter count, dataset size, and compute, given a sufficiently tuned learning rate. The Chinchilla follow-up [?] sharpened the compute-optimal trade-off between N and data tokens D . Both studies were conducted on natural-language corpora.

This project asks whether the same functional form holds on a markedly different distribution: SVG source code. SVG is structured, hierarchical, and heavily token-repetitive (tag names, attribute keys, numeric coordinates), so the entropy floor c in $L(N) = a \cdot N^{-\alpha} + c$ may differ substantially from prose, and the exponent α may shift if the redundancy structure changes the marginal value of additional capacity.

We pair this with a μ P sweep [?] to test the practical claim that the optimal learning rate transfers across width, which would remove the per-size LR sweep from the workflow at larger scales.

Contributions.

1. A reproducible 100M-token SVG training corpus built from `starvector/svg-stack` [?] with a 4096-token byte-level BPE tokenizer.
2. Fitted $L = a \cdot N^{-\alpha} + c$ scaling curves under SP for five model sizes (1M / 3M / 12M / 34M / 88M non-embedding parameters).
3. A side-by-side comparison of SP and μ P scaling, plus a $10\times$ extrapolation of the μ P-transferred LR.
4. Sample quality evaluation for the best model: perplexity, XML well-formedness rate, and `cairosvg` render rate, on unconditional and prefix-conditioned generations.

2 Data Collection and Preprocessing

2.1 Source dataset

The project specification recommends `starvector/svg-icons-simple` as the primary corpus, with optional supplementation from `svg-emoji-simple` or `svg-fonts-simple` to reach the 100M-token target. We instead use `starvector/svg-stack` [?] as a single source. The motivation:

- **Token budget without supplementation.** `svg-icons-simple` is 153 MB across ~ 89 K icons. After BPE tokenisation with vocabulary 4096 this yields well below the 100M training-token target, so the spec suggests combining multiple datasets. `svg-stack` alone reaches the budget, so all training data comes from one homogeneous distribution rather than a mixture whose composition would itself be a hyperparameter.
- **Distributional diversity.** `svg-stack` contains a broader range of SVG complexity (icons, diagrams, illustrations, decorative shapes) than the simplified-icons subset. This makes the validity metrics in Section ?? a stricter test of the model’s ability to generate well-formed SVG — well-formedness on simple icons is much easier than on the full stack distribution.
- **Pipeline correctness.** Our preprocessing applies the same coordinate-precision rounding, XML validation, and length filtering that the simplified variants receive upstream, so we can effectively re-create a comparably-clean corpus from the raw stack.

The SVG content is in the column `Svg` (capital S, verified via `ds.column_names`). We acknowledge this is a deviation from the recommended dataset and report all metrics on this corpus.

2.2 Cleaning and normalisation

Each SVG string is passed through three normalisation steps in `data/prepare.py`:

1. **Strip XML comments** via the regex `<!--.*?-->` with `re.DOTALL`.
2. **Collapse whitespace** (any run of spaces, tabs, or newlines becomes a single space).
3. **Round numeric values** to one decimal place. SVG path commands and attribute values typically carry several decimals of precision (e.g. `12.34567`). Rounding collapses near-duplicate numeric strings into a single BPE token, materially reducing vocabulary fragmentation. The precision is exposed as `coord.precision` in `configs/base.yaml`.

After cleaning, we apply two filters: a length filter (`min_chars=50` per the project spec, `max_chars=8000` to target ≤ 1024 BPE tokens at typical SVG density) and an XML well-formedness check via `lxml.etree.fromstring`. Documents that fail to parse are dropped. Render-validity via `cairosvg` is checked on a sample at evaluation time rather than per document during preprocessing, since rendering all documents would be prohibitively slow.

2.3 Tokenizer

A byte-level BPE tokenizer (HuggingFace `tokenizers`) is trained on the *cleaned* train split with vocabulary size $V = 4096$ and a single special token `<endoftext>`. Vocabulary size was chosen to balance two pressures specific to SVG: heavy repetition of tag and attribute substrings rewards larger merges, while the small-model regime (~ 1 M parameters) penalises wasting capacity on a

large embedding matrix. Within the 1K–8K range called for by the project spec, $V = 4096$ sits at a reasonable middle.

The trained `vocab.json` and `merges.txt` are committed to the repository so all downstream scripts share the exact same tokenisation across runs (BPE training order can be sensitive to multi-process parallelism).

2.4 Splits and tokenisation output

We use a 98/1/1 train/validation/test split by *file count*, seeded with 42. Splitting by file keeps every SVG intact in a single split, which is necessary for the XML-validity and render-validity metrics in Section ???. Each split is tokenised with the trained BPE; SVGs whose token length exceeds $T_{\max} = 1024$ are dropped (rather than truncated, which would silently bias the corpus toward partial SVGs). Token streams are EOT-separated and saved as flat `uint16` arrays (`train.npy`, `val.npy`, `test.npy`); `uint16` is the smallest dtype that holds $V = 4096$, halving disk and memory bandwidth versus `uint32`. The training array is hard-capped at `token_budget = 108`.

2.5 Statistics and examples

Figure ?? shows the distribution of document lengths in BPE tokens. The vertical dashed line marks $T_{\max} = 1024$, the cutoff applied to drop overlong SVGs. Figure ?? shows rendered SVGs sampled at the 10th, 30th, 50th, 70th, and 90th token-length percentiles, illustrating the complexity range that the model is trained on.

3 Transformer Scaling Study (SP)

3.1 Architecture

We adopt the nanoGPT implementation [?] directly, which is itself a clean re-implementation of GPT-2 [?]. The deviations from the base nanoGPT code are minimal and listed at the end of this subsection. All equations below correspond one-to-one with the nanoGPT source.

Forward pass. Token indices $\mathbf{x} \in \mathbb{Z}^T$ and position indices $\mathbf{p} = [0, \dots, T-1]$ are each looked up in learned embedding tables $W_e \in \mathbb{R}^{V \times C}$ and $W_p \in \mathbb{R}^{T_{\max} \times C}$:

$$\mathbf{h}_0 = W_e[\mathbf{x}] + W_p[\mathbf{p}],$$

followed by L transformer blocks and a final layer normalisation before the output projection:

$$\mathbf{h}_\ell = \text{Block}(\mathbf{h}_{\ell-1}), \quad \ell = 1, \dots, L, \quad \text{logits} = \text{LN}(\mathbf{h}_L) W_e^\top.$$

The output projection reuses $W_e^\top$ (*weight tying* [?]), coupling the token embedding and unembedding representations.

Transformer block (pre-LN). Each block applies causal self-attention and an MLP with pre-LayerNorm [?] and residual connections:

$$\mathbf{h} \leftarrow \mathbf{h} + \text{Attn}(\text{LN}(\mathbf{h})), \tag{1}$$

$$\mathbf{h} \leftarrow \mathbf{h} + \text{MLP}(\text{LN}(\mathbf{h})). \tag{2}$$

figures/seq_len_histogram.pdf

Figure 1: Document-length distribution in BPE tokens (training split). The dashed line marks the $T_{\max} = 1024$ cutoff used to drop overlong SVGs.

Causal multi-head self-attention. A single fused linear $W_{qkv} \in \mathbb{R}^{C \times 3C}$ produces queries, keys, and values, which are split and reshaped into h heads of dimension $d = C/h$. Scaled dot-product attention is computed per head:

$$\text{Attn}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}} + M\right) V,$$

where $M_{ij} = -\infty$ for $j > i$ and 0 otherwise enforces causality. We use PyTorch’s `scaled_dot_product_attention` (Flash Attention [?]) when available. This is SP; $1/\sqrt{d}$ is replaced by $1/d$ under μP (Section ??).

Feed-forward network. Two linear layers with $4\times$ hidden expansion and GELU:

$$\text{MLP}(\mathbf{h}) = \text{GELU}(\mathbf{h} W_1) W_2, \quad W_1 \in \mathbb{R}^{C \times 4C}, W_2 \in \mathbb{R}^{4C \times C}.$$

Initialisation. All weights are initialised $\mathcal{N}(0, 0.02^2)$; biases to zero. Residual projection weights (W_o and W_2) receive an additional scaling of $1/\sqrt{2L}$ per the GPT-2 paper [?], so that the residual stream variance stays bounded at initialisation regardless of depth.

Parameter count. The scaling law is fit against *non-embedding* parameter count N . Following nanoGPT exactly, only the position embedding is subtracted; the token embedding is retained because weight tying means those parameters directly participate in the output projection:

$$N = N_{\text{total}} - \underbrace{T_{\text{max}} \cdot C}_{\text{pos. emb}}.$$

Adaptations from nanoGPT. Our implementation differs in three ways: (1) vocabulary size $V = 4096$ instead of 50 257; (2) all biases disabled (`bias=False`), which is marginally better and cleaner for scaling analysis; (3) the `block.size` parameter is renamed `max_seq_len` for clarity. Width and depth are varied jointly across the five configurations to keep the aspect ratio L/C approximately constant.

What was borrowed vs. implemented. Per the project specification we explicitly itemise our use of nanoGPT [?]:

Component	Provenance
<code>CausalSelfAttention</code> , MLP, Block, GPT class structure	Adopted from nanoGPT verbatim. Same module hierarchy, same parameter naming (<code>c_attn</code> , <code>c_proj</code> , <code>c_fc</code>).
Pre-LN block, weight tying, scaled init for residual projections	Adopted from nanoGPT.
Flash Attention via	Adopted from nanoGPT.
<code>F.scaled_dot_product_attention</code>	
<code>get_num_params(non_embedding=True)</code> formula	Adopted from nanoGPT (subtracts only <code>wpe</code> , since <code>wte</code> is tied to the output and therefore active).
μP attention scaling, <code>set_base_shapes()</code> wiring, MuAdamW	Implemented by us, on top of the nanoGPT skeleton. The pre-scale-q-by- $1/\sqrt{d}$ trick (so Flash Attention’s internal $1/\sqrt{d}$ produces an end-to-end $1/d$ factor) is novel to this implementation.
Token-budget driven training loop, gradient accumulation, 1-epoch step computation	Implemented by us. The shape mirrors nanoGPT’s <code>train.py</code> but expects a token-budget config rather than a fixed step count.
LR sweep harness with <code>--mup</code> flag, sample-quality evaluator, sequence-length and rendering analysis	Implemented by us.

Table 1: Components borrowed from nanoGPT vs. implemented in this project.

3.2 Configurations

Widths and depths are chosen so that the non-embedding parameter count $N = C(V+1+2L) + 8LC^2$ (with $V=4096$) hits five logarithmically spaced targets. Each configuration uses MLP expansion factor 4 (i.e. $d_{\text{ff}} = 4C$) and head dimension $d_{\text{head}} = C/h \in \{32, 64\}$ to keep heads a power of two.

Config	L	h	C	d_{head}	$N_{\text{non-emb}}$
1m	4	4	128	32	1,049,728
3m	4	8	256	32	3,148,032
12m	8	6	384	64	11,016,576
34m	16	8	512	64	35,668,480
88m	18	12	768	64	88,108,800

Table 2: Model configurations. L = layers, h = heads, C = embedding dim, $d_{\text{head}} = C/h$. $N_{\text{non-emb}}$ computed analytically and verified at runtime via `get_num_params()`.

The PDF table suggests slightly different layer counts at the medium-to-large sizes (e.g. 6 layers / $C=384$ for $\sim 10\text{M}$; 10 layers / $C=512$ for $\sim 30\text{M}$). We adjusted depth so that each configuration sits cleanly on a target $N_{\text{non-emb}}$ multiple of $\sim 3\times$ the previous, which gives a more uniform spacing on the log- N axis for the power-law fit. Aspect ratios L/C remain comparable across the family (range 0.016–0.031).

3.3 Training Protocol

Optimizer. AdamW [?] with $\beta_1=0.9$, $\beta_2=0.95$, gradient clip 1.0. Weight decay 0.1 is applied only to parameters with $\text{ndim} \geq 2$ (weight matrices and embeddings); one-dimensional parameters (LayerNorm scales) are excluded, following the nanoGPT convention. Fused AdamW is used on CUDA when available.

Learning rate schedule. Cosine decay with linear warmup over the first 5% of steps:

$$\eta(t) = \begin{cases} \eta_{\max} \cdot \frac{t}{t_{\text{warm}}} & t < t_{\text{warm}}, \\ \eta_{\min} + \frac{1}{2}(\eta_{\max} - \eta_{\min}) \left(1 + \cos\left(\pi \frac{t - t_{\text{warm}}}{T - t_{\text{warm}}}\right) \right) & t \geq t_{\text{warm}}, \end{cases}$$

with $\eta_{\max} = 3 \times 10^{-4}$ and $\eta_{\min} = 3 \times 10^{-5}$ as defaults, swept per model size.

Batch size and gradient accumulation. The effective batch size is fixed at $B_{\text{eff}} = 524,288$ tokens for all model sizes. Since smaller models can accommodate only a modest number of sequences in a single forward pass, we use gradient accumulation:

$$k = \frac{B_{\text{eff}}}{T_{\text{max}} \times B_{\mu}} = \frac{524,288}{1024 \times 8} = 64 \text{ micro-steps},$$

where $B_{\mu} = 8$ is the micro-batch size. The loss is divided by k before each backward pass so gradients average (not sum) across micro-batches.

Training duration. Per the project specification, every model is trained for exactly one epoch over the 100M-token corpus. With $B_{\text{eff}} = 524,288$, this is $\lceil 10^8 / 524,288 \rceil = 191$ optimiser steps. The actual step count is computed at runtime by `train.py` when `max_steps` is unset (`null`) in the config: see the `resolve_max_steps` helper. The same step count is used for all sizes (each on the same data), keeping the training-token budget exactly matched across the scaling sweep.

Tracked metrics. Each step’s CSV log records: train loss, val loss (every `eval_interval` steps), learning rate, wall-clock seconds, tokens per second (computed as $B_{\text{eff}}/\text{step_time}$), and peak GPU memory in GB (`torch.cuda.max_memory_allocated` on CUDA, 0 on MPS). The best checkpoint by validation loss is saved as `best.pt` during training; periodic checkpoints are off by default (`save_interval` is set very high in `configs/base.yaml`).

3.4 Learning Rate Sweep

Per the project specification, the LR sweep is performed once on the *smallest* model (1m) and the winning LR is then re-used for all five SP sizes without retuning. We sweep five candidates $\{10^{-4}, 3 \times 10^{-4}, 10^{-3}, 3 \times 10^{-3}, 10^{-2}\}$ on 15% of the full one-epoch budget. For each candidate η , the minimum LR is set to $\eta/10$ and the model is re-initialised from scratch to eliminate order effects. The lowest-val-loss LR is committed to the per-size configs.

Figure ?? shows the sweep result. The winning LR is indicated; this same value is used for all five SP runs in Section ??.

3.5 Power-Law Fit

We fit $L(N) = a \cdot N^{-\alpha} + c$ to the five (N, L^*) points via `scipy.optimize.curve_fit` with positive bounds on a, α, c and report 95% confidence intervals from the covariance matrix. Figure ?? shows the training-curves panels (SP and μ P side-by-side); Figure ?? shows the fitted scaling laws.

Compute footprint. Training throughput and peak GPU memory per size are recorded in each run’s `log.csv` and summarised in Figure ??. The smaller models are bandwidth-bound; the larger ones become memory-bound on a single L4.

Results. *[Filled in once Colab runs complete; the analysis notebook regenerates Figs. ??–?? from the per-run CSVs and writes them to `report/figures/`.]*

4 μ P Scaling and Extrapolation

4.1 Reparameterization

We use Microsoft’s `mup` package [?]. Three changes from the SP implementation:

Attention scaling. The $1/\sqrt{d}$ factor in SP is replaced by $1/d$. Because PyTorch’s Flash Attention kernel applies $1/\sqrt{d}$ internally, we pre-scale the query by an additional $1/\sqrt{d}$ before the kernel call, so the net effect is $1/d$:

$$\tilde{Q} = Q/\sqrt{d}, \quad \text{Attn}(\tilde{Q}, K, V) = \text{softmax}\left(\frac{\tilde{Q}K^\top}{\sqrt{d}}\right) V = \text{softmax}\left(\frac{QK^\top}{d}\right) V.$$

The manual fallback path sets the scale to $1/d$ directly.

Base shapes and optimizer. A proxy base model with fixed width $C_0 = 128$ (same depth and head count as the target model) is instantiated and passed to `set_base_shapes(model, base_model)`. This attaches an `infshape` attribute to every parameter, which `MuAdamW` uses to apply the correct per-parameter LR multiplier. The base model is deleted immediately after to free memory.

Initialization-order invariant. `set_base_shapes()` *must* be called before the optimizer is constructed. If the optimizer is built first, `MuAdamW` reads `infshape` at construction time; without it the per-parameter scaling is silently absent and the run behaves identically to SP. In `mup_train.py` this ordering is enforced by the `configure_mup_model()` function, which calls `set_base_shapes()` and returns the model before `configure_optimizer()` is ever called.

The base width $C_0 = 128$ is held constant across all five model sizes so that the LR scaling ratios are consistent within the sweep.

4.2 Sweep and Comparison

We re-run the same five sizes under μP , with the LR *not* re-tuned per size — only the 1m μP model is swept (using `python sweep_lr.py --mup ...`), and the winning LR is transferred unchanged to 3m / 12m / 34m / 88m. The SP and μP 1m sweeps share the same five candidates and the same protocol. Both fitted scaling curves are overlaid in Figure ??.

4.3 Extrapolation

Using the better of the two fitted curves (likely μP , by design), we extrapolate the predicted validation loss to $N_{\text{target}} = 10 \times N_{88\text{m}} \approx 8.8 \times 10^8$ parameters. Confidence intervals on the prediction come from the covariance of the power-law fit, propagated analytically through $L = a \cdot N^{-\alpha} + c$. We do *not* train a model at this scale — it is well outside our compute budget on Colab L4 — so the extrapolation is an unverified prediction. The discussion below addresses the trustworthiness of this prediction.

Trust region of the extrapolation. The fit interpolates over ~ 1.7 decades in N (1 M to 88 M); extrapolating to 880 M is one further decade beyond this. Power laws empirically hold over several decades for natural-language LMs [?, ?], but several factors might cause the SVG curve to deviate:

- **Token budget mismatch.** At $N = 880\text{M}$ trained on the same 100M-token corpus we are nowhere near the compute-optimal token-to-parameter ratio [?] (Chinchilla suggests ~ 20 tokens/parameter). The extrapolation assumes the data side is held fixed; in practice the loss at $N = 880\text{M}$ would be data-bottlenecked, not capacity-bottlenecked.
- **Domain ceiling.** SVG has a finite irreducible entropy — the floor c in the power-law fit. If c is mis-estimated (e.g. because the small models haven’t gotten close enough to c), the predicted loss at large N has high variance.

Results. *[Filled in by the analysis notebook; extrapolated L_{pred} and predicted perplexity printed by the cell labelled “ μP extrapolation”.]*

5 Best Model Training and Sample Generation

We choose the largest model that fits our compute budget (88m) and the better of the two parameterisations (typically μP , but determined by the empirical scaling-law fit). Quantitative metrics reported:

- **Perplexity** on the held-out test split.
- **XML well-formedness rate:** fraction of generated samples that parse via `lxml.etree`.

- **Structural validity rate:** subset of XML-valid samples that have an `<svg>` root and at least one of `viewBox` or `(width, height)` attributes.
- **Render rate:** fraction (overall, and within XML-valid) that successfully render to PNG via `cairosvg`.

Qualitative analysis includes (i) at least 10 unconditional samples generated from a single `<endoftext—¿—` seed; (ii) at least 5 prefix-conditioned completions with prompts like `<svg viewBox="0 0 24 24"><circle cx="8"` (partial face), `<path d="M 10 10 L 50 10 L 50 50` (open path), and `<g transform="translate(10,10)"><rec` (group with one shape).

Evaluation procedure. Perplexity is computed over the held-out test split as token-level cross-entropy on non-overlapping windows of length $T_{\max} = 1024$, exponentiated. All windows in the test split are used unless noted. For sample-based metrics we draw K unconditional samples seeded with a single `<endoftext—¿—` token, attempt to parse each with `lxml.etree.fromstring`, and report the well-formedness rate. The XML-valid subset is then passed through `cairosvg.svg2png`; we report render rate both within the valid subset and as a fraction of all K samples. `cairosvg` failures (e.g. from unsupported features or malformed paths) are caught and counted as render failures rather than re-raised.

Sampling procedure. Generation is autoregressive with both temperature scaling and a composable top- k / top- p (nucleus) filter:

1. Compute logits at the last position; divide by temperature τ .
2. If `top_k` is set, set all logits below the k -th largest to $-\infty$.
3. If `top_p` is set, sort the surviving logits by probability, compute the cumulative softmax, and mask all tokens whose cumulative mass exceeds p (always keeping at least the top-1 token).
4. Softmax and sample via `torch.multinomial`.

Sequences exceeding $T_{\max} = 1024$ are left-cropped so the model always sees the most recent context. Sampling halts when an `<endoftext—¿—` token is emitted; the EOT itself is trimmed before decoding so output strings are clean SVG. Unconditional generation seeds with a single EOT token (matching the EOT separators used during training); prefix-conditioned generation tokenises the user prompt and uses those ids as the seed.

We sweep three temperatures $\tau \in \{0.5, 0.8, 1.0\}$ for the quantitative metrics (`evaluate.py --temperature_sweep 0.5 0.8 1.0`); the qualitative grids in Figure ?? use $\tau = 0.8$ and top- $p = 0.9$, which we found to balance coherence and diversity.

Results. *[Filled in once Colab runs complete; quantitative metrics from `results/best_eval.json`, qualitative grid from `generate.py --output_dir samples/best/`.]*

6 Design Decisions and Analysis

All quantitative analysis is reproduced by `analysis.ipynb`, which consumes the per-run logs (`checkpoints/{size}/log.csv`), LR-sweep results (`sweep_results.csv`), and evaluation outputs (`results/{size}_eval.json`) emitted by the training and evaluation scripts. The notebook produces every figure cited below and writes them to `report/figures/`.

Per-run training metrics. For each model size and parameterisation we plot validation loss against training step, with the LR schedule overlaid. These curves verify that each run actually converges within its step budget and that the LR schedule (cosine with 5% warmup) is taking effect.

LR sweep visualisation. For each size we plot validation loss against learning rate ($\log x$). Under SP we expect the optimum to drift with width; under μ P we expect the curves to align so the same LR is near-optimal for every size — the empirical signature of successful zero-shot LR transfer.

Scaling-law fit. The headline figure. We plot (N, L^*) on log-log axes for both SP and μ P and overlay the fitted curves $L = a \cdot N^{-\alpha} + c$. The fitted exponent α is compared to the Kaplan et al. value of $\alpha \approx 0.076$ for natural language; a substantial difference would be evidence that SVG’s repetition structure shifts the scaling behaviour relative to prose. The fitted entropy floor c is compared to the empirical per-token entropy of the test split.

μ P extrapolation. Using the μ P fit, we predict L^* at $\sim 10\times$ the largest fitted N (~ 880 M parameters). If a held-out run at that scale is trainable within the project’s compute budget, we plot it as a single measured point against the predicted curve.

Sample-quality scaling. A grouped bar chart of XML well-formedness rate and `cairosvg` render rate per size, alongside a log-log plot of test perplexity vs. N . The question is whether sample validity tracks scale smoothly or shows discrete capability jumps.

Discussion: design decisions and their impact.

- **Tokenisation.** BPE with $V = 4096$ was chosen at the lower-middle of the spec-recommended 1K–8K range. SVG is heavy on repeated substrings (`<path d=`, `fill=`#, common attribute keys), so a moderate vocab captures these patterns while keeping the embedding matrix small enough not to dominate the 1M model’s parameter budget. Coordinate-precision rounding (Section ??) removes a major source of vocabulary fragmentation; without it the BPE wastes merges on near-duplicate numerics.
- **Architecture choices.** Decoder-only Transformer with pre-LN, weight tying, and Flash Attention (Section ??). We deviate from the spec-suggested layer counts to space configurations more uniformly on the log- N axis, which gives a tighter power-law fit. Aspect ratios L/C remain in a narrow range so depth-vs-width is not a confound in the scaling fit.
- **Training decisions.** AdamW $(\beta_1, \beta_2) = (0.9, 0.95)$, weight decay 0.1 on 2D params only, gradient clip 1.0, cosine LR with 5% warmup, bfloat16 on CUDA. One epoch over the 100M-token corpus matches the spec’s “1 epoch for the scaling plot” requirement exactly. Effective batch size $B_{\text{eff}} = 524,288$ tokens is held constant across all sizes, so per-step compute scales only with N .
- **Scaling insights.** *[Empirical: compare fitted α to Kaplan et al. $\alpha \approx 0.076$. SVG’s syntactic constraints could push α either way — higher if the redundant structure is reducible by capacity, lower if there is a hard syntactic ceiling. The fitted c should approach the empirical per-token entropy of the test split.]*

- **Learning-rate scaling: SP vs μ P.** *[Empirical: the SP curve (LR fixed at the 1m-tuned value) is expected to under-perform at large N as that LR becomes too large. μ P with the same 1m-tuned LR should track or exceed the SP curve at large N. The exact crossover size is the practical takeaway.]*
- **SVG-specific patterns.** *[Empirical: at small N we expect the model to learn local syntax (matched tags, valid attribute names) without coherent geometry. At large N we expect spatial-coherence emergence: bounded shapes within `viewBox`, sensible path-closing, common color palettes. Phase-transition vs smooth-improvement is an empirical question.]*
- **Challenges encountered.** *[Empirical: what didn't work, what we'd do differently with more compute — e.g. the 880M extrapolation point we couldn't train, the μ P coord-check we did not run.]*

7 Conclusion

We trained five decoder-only Transformer language models from ~ 1 M to ~ 88 M non-embedding parameters on a 100M-token corpus of SVG source code, fitted a power-law scaling curve under both standard parameterisation and μ P, and used the better fit to extrapolate predicted loss at $10\times$ the largest fitted scale. We additionally trained the largest model to convergence, generated samples under multiple sampling strategies, and measured XML well-formedness, structural validity, and render rate as quantitative measures of generation quality.

[Empirical-summary paragraph: report fitted α and c for SP and μ P, the comparison to natural-language scaling, the extrapolation prediction with confidence interval, and the best model's perplexity / sample-quality numbers.]

A Generated SVG sources

[Source listings of the rendered samples in Figure ??, plus the prefix-conditioned completions, will be appended here from `samples/best/sample_*.svg`.]

References

- [1] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, I. Sutskever. *Language Models are Unsupervised Multitask Learners*, 2019. <https://openai.com/research/language-unsupervised>.
- [2] O. Press and L. Wolf. *Using the Output Embedding to Improve Language Models*, EACL 2017. <https://arxiv.org/abs/1608.05859>.
- [3] T. Dao, D. Fu, S. Ermon, A. Rudra, C. Ré. *FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness*, NeurIPS 2022. <https://arxiv.org/abs/2205.14135>.
- [4] A. Karpathy. *nanoGPT*, 2022. <https://github.com/karpathy/nanoGPT>.
- [5] R. Xiong, Y. Yang, D. He, et al. *On Layer Normalization in the Transformer Architecture*, ICML 2020. <https://arxiv.org/abs/2002.04745>.
- [6] J. Kaplan, S. McCandlish, T. Henighan, T. B. Brown, et al. *Scaling Laws for Neural Language Models*, 2020. <https://arxiv.org/abs/2001.08361>.

- [7] J. Hoffmann, S. Borgeaud, A. Mensch, et al. *Training Compute-Optimal Large Language Models*, 2022. <https://arxiv.org/abs/2203.15556>.
- [8] G. Yang, E. J. Hu, et al. *Tensor Programs V: Tuning Large Neural Networks via Zero-Shot Hyperparameter Transfer*, 2022. <https://arxiv.org/abs/2203.03466>.
- [9] I. Loshchilov and F. Hutter. *Decoupled Weight Decay Regularization*, ICLR 2019. <https://arxiv.org/abs/1711.05101>.
- [10] J. A. Rodriguez, S. Agarwal, I. H. Laradji, et al. *StarVector: Generating Scalable Vector Graphics Code from Images*, 2023. <https://arxiv.org/abs/2312.11556>.




Figure 2: Rendered SVG examples from the training split, sampled across the length distribution (10th, 30th, 50th, 70th, 90th percentile).





Figure 3: LR sweep on the 1m model: best validation loss vs. candidate LR. Per the spec, the winning LR is re-used unchanged for all larger SP sizes.



figures/training_curves.pdf

Figure 4: Validation loss vs. training step for each model size. **Left:** standard parameterisation. **Right:** μ P. Same y-axis for visual comparability.





Figure 5: Scaling law fit $L = a \cdot N^{-\alpha} + c$ on log-log axes. SP and μ P measured points and fitted curves overlaid. Fitted exponents and entropy floors annotated in the legend.



figures/compute_footprint.pdf

Figure 6: Tokens/sec (left) and peak GPU memory (right) vs. non-embedding parameter count, on Colab L4 with bfloat16. Both axes log-scaled where appropriate.

figures/sample_quality.pdf

Figure 7: **Left:** XML well-formedness and render rate vs. model size. **Right:** test-set perplexity vs. non-embedding parameter count (log-log).



Figure 8: Rendered grid of unconditional samples from the best checkpoint ($\tau = 0.8$, top- $p = 0.9$). The full SVG source for each is reproduced in Appendix ??.